

Project 1 reviews have been re-issued — here is where everyone stands

01-ml-project — Customer Segmentation · 9 September 2026

The Monday catch-up deadline has passed. I have re-read every group that pushed something and rewritten the reviews against what actually arrived. **Ten of fifteen A1 groups are now accepted.**

Pull and re-read your review

```
git pull origin main
```

Then open `submissions/a1-mon-thu/<your-group-id>/01-ml-project/<your-group-id>_final-review.md`

If your outcome changed, the file says so at the top and is dated **09/09/2026 (round 2)**. The 04/09 version no longer applies.

Where each group landed

GROUP	ON 04/09	NOW	WHAT HAPPENED
a1-g02	Revise	Accepted	All four fixes. Personas <i>derived</i> from the table, not renamed.
a1-g03	Revise	Accepted	Outputs saved — that was the blocker.
a1-g10	Revise	Accepted	All three fixes plus both notes. Strongest turnaround in the cohort.
a1-g13	Not submitted	Accepted	Delivered on the deadline. Best k justification in either cohort.
a1-g04	Revise	Revise	Pushed on 06/09, but neither required fix.
a1-g06	Revise	Revise	Nothing pushed.
a1-g08	Revise	Revise	Nothing pushed.

Already closed out and needing nothing further: a1-g01, a1-g05, a1-g07, a1-g09, a1-g11, a1-g12, a1-g14, a1-g15.

Final catch-up deadline — Sunday 13/09, end of day. This applies to **a1-g04, a1-g06 and a1-g08** only, and it is the last extension on Project 1. After Sunday the outcome is recorded as it stands.

If something is blocking you — git, the dataset path, a notebook that errors when you actually run it, anything inside the group — message me today. Every one of those is a fifteen-minute conversation. I cannot help with silence, and for two of you this is now the second deadline it has cost.

Thursday 10/09 — half discussion, half new material

We are **still at the end of Session 10**. Monday closed S10 out with the discussion and the mini project, so Thursday is the first half of **S11**, with the first half of the meeting given to the last Project 1 discussion round.

DATE	WHAT IT CARRIES	PRESENTING
Thu 10/09	½ discussion + first half of S11	a1-g02, a1-g10, a1-g13

Mon 14/09	Second half of S11 + short catch-up slot	a1-g04, a1-g06, a1-g08 (if delivered)
Thu 17/09	S12	—

Thursday is only half a meeting, so the three slots are **five minutes each on one specific thing** — not a walkthrough:

- **a1-g02** — why cluster 1 is “Standard / Mainstream” and not “VIP”.
- **a1-g10** — the k=2 and k=3 profiles side by side, and what changed about the business recommendation.
- **a1-g13** — the k justification only: why they overrode the silhouette peak.

Three things worth everyone reading

Naming the boring majority correctly is the hardest part of this project. Four groups across both cohorts had a cluster of roughly 240–260 ordinary customers — a couple of months since the last purchase, eight or nine purchases each — and called it “VIP”, “At-Risk” or “Dormant”. Every one was returned for it. That cluster is your business. If your persona name would send the marketing team to spend win-back budget on 80% of the customer base, the name is wrong even when the clustering is right.

Write down the trade-off you actually made. The strongest k justifications this round were the ones that conceded something. a1-g13 wrote that silhouette preferred k=2 and that their k=4 scored *lower* — and then argued for k=4 anyway, on grounds they could defend. That is harder to write than a clean answer and worth more. Describing the evidence as agreeing with you when it does not is the one thing a reviewer always catches.

`git add .` from the repo root is how files end up where they should not be. This round produced a stray `FETCH_HEAD` at the repo root, a duplicate notebook committed beside the original instead of replacing it, and several copies of the shared CSV inside group folders. Use `git add <your folder>`, and read `git status` before you commit. Nothing outside your own group folder should ever appear in it.

On session numbers

The schedule files now key everything on **dates** rather than session numbers, because a session can span two meetings or share one with a discussion round, and the two had drifted apart. `discussion/a1-discussion-schedule.md` has a table showing what each date carries. A3 is now ahead of us on session numbers — the project deadlines stay shared, the session numbers do not.

Questions on your own review come to me directly — your mentor’s stage is finished. Anything about this announcement goes in the comments so the whole cohort sees the answer once.

— Hozaiifa Moustafa